

Paul Johnson , Director, CRMDA <pauljohn@ku.edu>

September 1, 2017

This document discusses the CRMDA package templates and how to use them. It compares some strengths and weaknesses of the ‘noweb’/L^AT_EX format versus ‘Rmarkdown’. It also mentions two code-chunk processing engines, ‘Sweave’ and ‘knitr’.

1 crmda package document templates

The **crmda** package provides templates for several types of documents using various front and backends.

1.1 Mix and match

The **frontends** are Rmarkdown (suffix *.Rmd) and noweb/L^AT_EX(suffix *.Rnw).¹

The **backend** (output) formats are PDF and HTML. The formats are “guide” and “report”.

Code chunks in noweb/L^AT_EXdocuments can be processed either by **Sweave** or **knitr**.

The mix-and-match of the possibilities in the format “*frontend2backend*-doctype-chunktype” results in 6 document types (so far).

1. Guides

- rmd2html-guide
- rnw2pdf-guide-knit
- rnw2pdf-guide-sweave

2. Reports

- rmd2pdf-report
- rnw2pdf-report-knit
- rnw2pdf-report-sweave

¹A **noweb** file is a L^AT_EXfile that has code chunks embedded in it. Noweb documents can be prepared with [LyX](#).

1.2 Guides versus reports

A **guide** is intended for students or other learners. A guide will generally include code and output excerpts. Preferably, the code examples will have line numbers and syntax highlighting. In most cases, a guide format is not format suitable for reports to clients or for journal articles.

A **report** is a more formal document. A report is suitable for a journal article or a technical report. A report has less (maybe no) code and almost never will it include “raw output” from a computer program. A report includes closer-to-publishable tables and figures.

The style for “report” documents is fairly well settled. We have a header with the title, the logo, and author information. On page 1, we have contact information, in the style of a stationary footer.

The style for “guide” documents is not well settled. It is one of the problems on our immediate agenda.

1.2.1 What about slides?

We are 90% finished with a template style for rmd2pdf slides (Beamer).

It is ironic that this exercise began in 2015 as a search for workable slide shows via Rmarkdown documents. That search for a fully workable method to generate slides with markdown is, at the current time, a failure, as no avenue has been found that matches the needs we have.

1.2.2 What about Word documents

Soon to appear will be “docx-report” and “docx-guide”.

1.3 Let us manage defaults for you

The **crmda** package has settings that work for our documents.

Because the number of formats among which **pandoc** can convert is truly immense, and the number of options for each format is immense, *there are too many details*. This document undergoes a transitions **Rmd** -> **md** -> **utf8.md** -> **tex** -> **PDF**. The transition from **utf8.md** to **L^AT_EX**, considered in isolation, is nearly overwhelming. When I’m producing the PDF for this document, the **pandoc** command is elaborate.

```
$ /usr/bin/pandoc +RTS -K512m -RTS crmda-reports_and_guides.utf8.md -to latex -
from markdown+autolink_bare_uris+ascii_identifiers+tex_math_single_backslash -table-
of-contents -toc-depth 2 -template theme/crmda-boilerplate.tex -highlight-style had-
dock -latex-engine pdflatex -listings
```

These can be customized, both by revising the document header and by adjusting the build script, **rmd2pdf.sh**. In the **CRMDA**, we’d rather keep the output consistent by fixing the style sheets and processor scripts rather than hand-coding individual documents. We are somewhat, but not entirely, democratic in selecting styles.

The Rstudio team, the authors of the packages **rmarkdown** and **knitr**, and the program **pandoc** have created a family of settings and default styles that work for their purposes. In the **crmda** package,

we have customized document templates and style sheets that override some of their decisions. The functions `rmd2html` and `rmd2pdf` are “wrappers” that apply our organizational strategy.

The conclusion should be the following. We believe the default configuration in our document packages should “just work”. However, if specialized adjustment is necessary, we allow it. Over the long run, however, beneficial specializations should be incorporated into the document defaults.

2 Usage overview

Here we will illustrate the process of initiating a guide document that is to be prepared with Rmarkdown and will have the HTML backend.

Start R in a folder where you would like to create a write-up and run

```
> initWriteup("rmd2pdf-report")
```

We could have specified an output directory, but we did not. The default creates a folder named `writeup/rmd2pdf-report`. It will be created in the R current working directory. That folder should have a copy of the template, some instructions, and a compiler script.

If I were running this inside an R folder within one of our projects, I would instead run

```
> initWriteup("rmd2pdf-report", dir = "../writeup")
```

because I want my writeups to collect in the writeup folder of the project.

The first order of business is to rename the subfolder from “rmd2pdf-report” to something relevant to a writeup. There may be several reports for this project, choose a name that will help you guess which writeup is in which folder. Be aware that `initWriteup` will erase the files in that folder, if you don’t rename it/them.

Next, we need to make sure the build programs exist and work in your computer. Rename the file `report-template.Rmd` to match your purpose. My file name is “`crmda.Rmd`”.

Test the build tools. Edit your document, insert your name or a title and save the document. Now try to compile it.

The file can be compiled in two ways.

1. Open an R session and use the function `rmd2pdf()` in the `crmda` package. I’ll compile my file `crmda.Rmd`.

```
> library(crmda)
> rmd2pdf("crmda.Rmd")
```

The `rmd2pdf` is a *wrapper* function. It does not do any real work, it just collects your file name and applies defaults that we have set. It calls the `render` and `pdf_document` functions in the `rmarkdown`. Because `knitr` is the only method for handling code chunks in Rmarkdown documents, also involves the `knitr` package.

The `rmd2pdf` allows the user to add arguments that are passed along to `render` and `pdf_document`. This is explained in the help page for `rmd2pdf`.

2. Run a shell script provided with the template.

It is not necessary to open R to compile the document. The output can also be generated by running a command line script that does the same work.

```
$ ./rmd2pdf.sh crmda.Rmd
```

The rendered output will be a pdf file. The configuration inside `rmd2pdf.sh` indicates that output is written in the current working directory.

If one uses Rstudio as a text editor for an Rmd file, one may not obtain satisfactory output by compiling the document with the pull down menus.

3 About the shell scripts

In July, 2017, I decided to change the way that the shell scripts are designed. Users who inspect them should see the script runs R and launches a function in the `crmda` package. As a result, there should not be any difference in the result from running the script `rmd2pdf.sh` and the function inside R named `rmd2pdf`.

The only difference that might arise is when there are additional arguments supplied in the script that are not present in `rmd2pdf`.

One difference might be whether or not a table of contents is requested. The default arguments specified in the script are simple, but they do include the table of contents (this was inserted in this script in June 2017). The important line in `rmd2pdf.sh` is this:

```
defaults="toc=TRUE, output_dir=\"$pwd\""
```

This turns on the table of contents feature and asks for output in the current working directory. We have not worked on command-line argument processing in `rmd2pdf.sh`, so if a user wants to change the arguments passed along, it will be necessary to edit that script file and items in the `defaults` character string.

A little lower in `rmd2pdf.sh`, one should see that the business-end of the project is a call to `Rscript` which supplies the file name, along with the `defaults` string, in a not-very-subtle way:

```
Rscript -e "library(crmda); rmd2pdf(\"$filename\", $defaults)"
```

The new arguments that users might insert in `defaults` should match the arguments that might be supplied to the `rmd2pdf` function in the `crmda` package. Hence, the help page for `rmd2pdf` should be consulted for details. Arguments that can be specified for the `render` and `pdf_document` (or `html_document` if we are creating HTML output) are spelled out. Any defaults specified in the script, or in the function call to `rmd2pdf` inside R, will override the other customizations that are place in the `crmda` package.

We hope that only minor customizations are needed. The document you are reading now is produced with a revised version `rmd2pdf.sh`. By default, the `rmd2pdf` function will pull the $\text{L}^{\text{A}}\text{T}_{\text{E}}\text{X}$ template for `pandoc` from the `crmda` package's install folder. That is to say, the `rmd2pdf.sh` script overrides the template file that exists in the package install folder and instead uses the one in the theme subdirectory of the current document directory.

While building this, I had some trouble compiling the document. I turned on detailed error reporting (`quiet=FALSE`). I also turned off the `clean` argument that erased the intermediate markdown files and asked to keep the \LaTeX intermediate files. Hence, my `defaults` string has

```
defaults="toc=TRUE, output_dir=\"\$pwd\", clean=FALSE, keep_tex=TRUE,
        quiet=FALSE, template=\"theme/crmda-boilerplate.tex\""
```

Because I don't clean and keep the \LaTeX intermediate file, the working directory has these "extra" files when the run is completed:

```
crmda.utf8.md
crmda.knit.md
crmda.tex
```

Some of these argument can be specified in the document header. For example, the current document header has `keep_tex: false`. The script setting `keep_tex=TRUE` overrides that.

The fact that settings can be specified several places—in the document header, or in the processing script, or in the functions that handle documents—is confusing. I don't think anybody can deny that.

Inside R, using the function `rmd2pdf`, I achieve the same result by inserting the arguments into the function call:

```
> rmd2html("guide-template.Rmd", toc=TRUE, output_dir="$pwd",
          clean=FALSE, quiet=FALSE,
          template="theme/crmda-boilerplate.tex",
          keep_tex=TRUE)
```

The R help page for `rmd2pdf` should be helpful (or will be, eventually).

4 Working with \LaTeX

Like so many other concepts and tools in this area, \TeX was a creation of Donald Knuth at Stanford. \TeX was the precursor to \LaTeX .

The CRMDA maintains a Web page about \LaTeX : <https://crmda.ku.edu/latex-help>

That page has basic guides and information about the KU dissertation template.

I prefer to work with \LaTeX documents, for a number of reasons. The quality of the PDF output is nicer, in my opinion, and more predictable. However, the major reason I prefer \LaTeX is that the Sweave option `split=TRUE` is allowed. That option creates separate `*.tex` output files, for each code chunk. The developers of the Rmarkdown documents framework disapprove of `split` and elected, consciously, not to implement it.

A couple of the questions not considered in our Web page are the following.

5 Choosing among formats

5.1 Which backend?

Should I end up with HTML or PDF? The answer depends on the intended audience/client. If a “paper” must be submitted, obviously choose PDF. If the document needs numbered equations, cross references, and “floating” tables and figures, choose PDF. If the document is intended for a Webpage, then HTML is the obvious choice (unless you simply want to convey a PDF document via the Web). Our HTML template includes a cascading style sheet feature that allows both color-highlighted sections and tabbed sub-sections.

There is concern about mathematics in HTML documents that deserves mention. For many years, the inability of Web pages to display equations was a major problem. Many tedious, ugly methods were developed. A more-or-less workable solution was developed, a framework called MathJax. MathJax allows inclusion of math markup in the page which—when the conditions are right—can be converted by the Web browser to look like equations.

If the user is offline, or if the MathJax server is not available, then the HTML document’s math will now display. If one wants to put math into a document, using HTML is inherently risky. If one needs to be 100% sure that math will display as intended, choose to create PDF documents.

5.2 Which Frontend? Write in LaTeX or Rmarkdown?

This will be the answer:

When choosing the frontend, consider the backend. Where you want to end up determines where you start.

While working on this document, I prepared an original version in Rmarkdown that was compiled into HTML. Because some features failed to compile, I changed the backend to PDF in the report style. As a result, several features that are unique to the HTML backend had to be removed. HTML offers access to some special document formatting features that are simply not available in PDF, and the converse is also true.

I believe the following are good conclusions:

1. If one intends to export as HTML, then markdown is, *by far*, the most reasonable choice for a frontend. Markdown was developed, first and foremost, as a simpler way to generate Web pages.
2. If one intends to export to PDF, then markdown or \LaTeX can be useful. But \LaTeX is probably better. \LaTeX is primarily intended for the creation of publication-quality documents in PDF format. Conversion from \LaTeX to HTML is less decidedly less unsatisfactory.

Since one can put much \LaTeX markup into a markdown document, perhaps the difference is not so great as it seems. The Rmarkdown compilation process (see Appendix 6.1) generates a \LaTeX file at an intermediate stage, so in some sense the same PDF result ought to be possible with Markdown or \LaTeX document preparation. However, in practice, we find differences in conveniences for authors.

I would summarize the situation with a poem:

Markdown documents intended for HTML allow some \LaTeX code.
Markdown documents intended for PDF allow more \LaTeX code.
Almost all HTML code is tolerated well in a Markdown document intended for HTML.
No HTML code is tolerated in a document intended for PDF.

The main point is that if one writes a markdown document, using special features intended for the backend, then it is generally not possible to, at the last minute, change the output format from HTML to PDF, or vice versa. HTML output has advantages in Web style features, while PDF documents have advantages in “on paper” presentations.

For novices, \LaTeX seems more difficult than markdown. Perhaps this is not such a big hurdle as it used to be because [LyX](#) is available. In my opinion, [LyX](#) makes preparing a \LaTeX document much easier than it is to prepare a similarly complicated markdown document.

Important caution about the HTML backend: Math is not incorporated in HTML in the same way as PDF.

Compiling a document into PDF uses a program like `pdflatex` to put the equations “in” the document. They are displayed in (more or less) the same way on various browsers and operating systems. The same is not true for math in HTML documents. Simply put, **math is not allowed in HTML**. We think it is allowed—our eyes tell us it is allowed—because we browse Web pages that show equations. However, this is an illusion achieved by extraordinary measures involving Javascript and third party servers. The beautifully formatted \LaTeX equation is not “embedded” in the HTML, it is instead delivered as code “available for rendering” in the Web browser. The HTML code is converted, via javascript and functions supplied interactively from the MathJax Web server.

Markdown to HTML allows all valid HTML markup, but Markdown to PDF does not allow all \LaTeX . This is a somewhat surprising difference. All HTML markup I’ve tried works well in an Rmarkdown document that aims to go into HTML. However, not all \LaTeX markdown is allowed in an Rmarkdown document going to PDF. And even less \LaTeX code works well if the intended backend is HTML. One cannot insert italics with \LaTeX in a document intended for HTML. For example, writing `\emph{italics}` or bold `\textbf{bold}` in the style of \LaTeX code will have no effect in an HTML document. However, if PDF output is used, then both of those \LaTeX codes work. As evidence, note I get *italics* and bold **bold** in this PDF document.

5.3 Should one prefer Sweave or knitr?

This question is meaningful only in `noweb`/ \LaTeX documents. In Rmarkdown, `knitr` is the only available method to process code chunks. In `noweb`, one can choose between `Sweave` and `knitr`. Perhaps that suggests that, if one must learn one set of chunk options, then `knitr` options are the right place to start (since they can be used in documents intended for HTML or PDF).

The chunk options allowed the original `Sweave` were `echo` (include code with output?), `eval` (run the chunk calculations?), `include` (display the chunk in the document?), `fig` (code generates a figure?) and `results` (output in \LaTeX is handled differently than raw \TeX). There are a few others, but that is most of the story.

`knitr` honored most of the `Sweave` options and then added many more (see [knitr code chunk options](#)).

One benefit of Rmarkdown with `knitr` is that it is possible to make documents about other programs (not just R). I've explored knitr to weave documents about BASH shell programming, for example.

6 Troubleshooting

The first step is understanding the trouble. The trouble stems from the fact that each document must be transformed through several stages to reach the final result. Understanding that, and learning about the problems that appear at each stage, can help with the troubleshooting strategies that we recommend.

6.1 Compilation Stages

Steps to compile documents break down into 2 phases.

1. Handle code chunks
2. render the resulting document.

A noweb file is converted from Rnw into PDF by a sequence of transitions.

1. `Rnw` -> `tex`. This is called “weaving” or “knitting”, depending on whether Sweave or knitr is the code processing engine. R finds the code in the Rnw file and inserts results into a new LaTeX file. The difference between weaving and knitting will be explained below.
2. `tex` -> `pdf`. The default is `pdflatex` for this step, but the alternative `xelatex` is growing in popularity because it more gracefully handles Unicode characters (utf8).

If the document is edited with [LyX](#), there is an implicit step 0,

0. `LyX` -> `Rnw`. In the LyX pull down menu system, this is represented by Export -> Sweave.

A markdown file is converted from Rmd to HTML by this sequence of transitions

1. `Rmd` -> `md`. The “knitting” process replaces code chunks by R input and output, converting the R markdown file into an ordinary markdown file. In my system, an “md” file is generated, and then a second “utf8.md” is generated to clean up the file encoding.
2. `md` -> `HTML`. Currently, most people use the program `pandoc` for this. A version of `pandoc` is distributed with Rstudio for the convenience of users. Linux users probably have `pandoc` available as standard system packages and the Rstudio version be removed.

The production of PDF from markdown, involves an additional transition.

1. `Rmd` -> `md`. Knitting converts code chunks into R input and output that is inserted into an md file created in the `pandoc` markdown style.
2. `md` -> `tex`. A \LaTeX file is created by `pandoc`. In the header of the md document, one can set a number of parameters to alter the \LaTeX generation process. For troubleshooting, “keep_tex: yes” to keep the `tex` file.

3. `tex -> pdf`. The default program for this has been `pdflatex`. It may be important to know that a \LaTeX document may need to be run through `pdflatex` several times because cross-references among pages and equations need to be made consistent.

The R packages `rmarkdown` and `knitr` orchestrate the process that builds the instructions to `pandoc`. In `rmarkdown`, the function `render` orchestrates all of the work. It calls chunk calculator and assigns work among the various conversion programs. The functions in the `crmda` package named `rmd2pdf` or `rmd2html` are “wrapper” functions that adjust settings sent to `render`.

6.2 Avoiding compilation trouble

The document production phase can fail at many steps (see Appendix 6.1). While editing a document, authors are well advised to heed the advice:

Compile early, compile often!

When a mistake is inserted, it is best to find it as soon as possible.

6.3 When debugging, check intermediate files

The compiler scripts may erase intermediate files. While debugging a document, we want to disable that clean-up step so that we can see what goes wrong. In an Rmd document that ends up in PDF, for example, we should be able to inspect an `md` file and a `tex` file. We can not only inspect those files, but we can also attempt to compile them in isolation so that we can see what is going wrong.

This document includes the header argument `keep_tex: true`, which means we save a copy of “reports_and_guides.tex”.

While developing this document, some of the problems with backend-specific code have come to the forefront. The HTML backend allows pleasant color-coded section headings which included in the PDF output. Tables that work well in PDF documents don’t work in HTML, and vice versa. Some HTML tables that are legal HTML don’t cooperate with `pandoc`.

It is also worth mentioning that error messages are not always informative. In fact, we sometimes don’t get error messages when we should. Rather, we simply receive bad output. While developing this document I noticed that when a user includes erroneous \LaTeX code in a markdown document, a flawed HTML output is generated without error or warning. On the other hand, changing the intended backend to PDF causes the compiler to fail and issue an error message. If one is exporting to HTML, then, a very careful proofreading of the output to check conversion of \LaTeX code into HTML is necessary.

7 What do we Really, Really Need?

7.1 Things we wish we could have in HTML output (that we can get in PDF output)

Numbered equations, easy cross references, numbered tables and figures

7.2 Things we wish we could have in PDF output (that we can get in HTML)

A splash of color, mainly. This is possible in PDF, but more difficult, at least on the surface.

7.3 Math

We are a Center focused on methodology. It is necessary to be able to write about math, preferably with a standard, uniform mathematical markup language, such as \LaTeX .

Many social scientists are not familiar with \LaTeX document preparation. That was a hurdle that kept many authors with Microsoft Word, even when they were frustrated with it. The difficulty of using \LaTeX was solved, to a significant extent, by LyX, an open source graphical interface. With LyX, or other editors that could generate \LaTeX output (such as Scientific Word, TexMacs, or Abiword), authors who were not computer programmers could learn enough \LaTeX to finish their projects.

7.4 Literate documents: include code and output

We need to be able to write about computer code. The “old fashioned” way is to copy/paste code and output into documents. That’s somewhat error prone and difficult to keep up-to-date.

Donald Knuth, a famous Stanford professor of computer science, proposed strategies to integrate the production of documents with the development of computer code. Rather than creating code in one file, and documentation in another, the idea was that the two parts of our work should be blended in a “literate programming” exercise.

The literate programming idea is more a general way of life than it is a particular document production strategy. It is, partly, aimed at programmers who don’t like to write instruction manuals. In the end, however, it may have more impact on non-programmers who need to prepare technical reports that include computer code examples.

In computer programming, one of the biggest impacts of literate programming is the proliferation of systems for preparation of documentation within code files. In the 1990s and early 2000s, when I was working on the Swarm Simulation System, we used a framework called Autodoc that allowed us to write instructions into Objective-C code that were later harvested and turned into instructional manuals. See, for example, the documentation for the [Opinion Formation model] (<http://pj.freefaculty.org/Swarm/MySwarmCode/OpinionFormation/Opinion-Docs>). The Autodoc program was poorly documented and not easy to get, but soon after that, a new coding system called Doxygen became widely available. Doxygen, developed for creation of instruction manuals for C++ programs, was a major success in computer programming. Like Autodoc, Doxygen gave programmers a relatively convenient method to explain what they were doing without wasting too much time.

In the modern experience of most CRMDA staffers, the “documentation inside code” approach is visible in the Roxygen markup method used for functions in R documents. Any function worth using should have Roxygen markup.

7.5 Where we have been

In order to embrace the importance of using either markdown or \LaTeX , one must first abandon the idea that Microsoft Word can ever be useful for serious authorship. That's a big step for many graduate students and professors.

If we look past GUI “what you see is what you get” (*wyswig*) word processors, where do we go? For a long time, the only answer was \LaTeX . However, there was a fatal weakness in \LaTeX . It is intended for PDF output, not Web pages. Exports into HTML were problematic. \LaTeX was not only difficult for some to use, but it also did not benefit from fancy features that were becoming available in the Internet, especially cascading style sheets and Javascript.

This gap in the document production process created a need for a new methods. In the 1990s, there was quite a bit of effort to make user friendly Web page editors, so that authors could have a Word-like experience that would generate HTML. The end result, generally, was difficult-to-maintain HTML documents. It was generally not feasible to edit and revise documents, the HTML generated was both extremely complicated and generally unsatisfactory.

Another strategy was the development of alternative markup languages. In a way similar to \LaTeX , these markup languages (e.g., *docbook*) drop the idea that the user should have a *wyswig* experience. Instead, the author would again become a programmer who would insert symbols to create sections.

Markdown was developed as a rejection of both ugly markup documents (either \LaTeX or HTML) and *wyswig* editors. The idea is that documents should be text files that are readable *as is* but also convertible into other backends. The “markdown” movement seeks to deliver an easier-to-edit, less difficult-to-read, and easier-to-convert format. The leader is John Gruber, whose Website is boldly named “[daringfireball](#)”.

Markdown is intended to be easy-to-read, so that even if it is not compiled into a backend, it might be presentable to an audience of non-programmers. That general idea seems unrealistic to me, there is almost never going to be an audience (for CRMDA, at least), which is eager to look at a markdown file. The dropback argument is that a markdown file is more easily produced by a novice who does not want to learn to use \LaTeX . This seems more persuasive.

The strength of markdown is that it makes it fairly easy to produce HTML documents that utilize some (not all) of the strengths of the World Wide Web's most commonly used method of communication. Where markdown is not capable, one can still write “raw HTML” in the middle of a markdown document.

8 Required footer

The guide documents we require authors include a final chunk that includes the session information, to be used in bug-tracking.

Reports do not include raw output, so it is not recommended to insert that raw output in the report. Instead, we ask report writers to include a final R chunk that saves the session information in a file in the same directory as the pdf output.

Available under [Created Commons license 3.0](#)